

✓ Predicting Student Success Using Logistic Regression

This notebook explores whether demographic attributes (gender, age, region, disability) can predict whether a student will pass or fail. The data is drawn from the Open University Learning Analytics Dataset (OULAD) and focuses on logistic regression classification using scikit-learn.

Project Objective

To test whether we can predict academic success (pass/fail) using demographic features available at the start of a course — and assess the predictive power of logistic regression models.

Dataset Overview

- **Source:** Public dataset from [OULAD on Kaggle](#)
 - **Tables used:**
 - `studentInfo`: student demographics and final outcomes
 - **Key Features:**
 - `gender`
 - `age_band`
 - `region`
 - `disability`
 - `final_result` (target, filtered to Pass or Fail)
-

Methodology

1. **Data extraction** from SQLite database
 2. **Data cleaning** and encoding of categorical variables
 3. **Logistic regression model** using:
 - Basic features (unweighted)
 - Balanced class weights to address target imbalance
 4. **Performance evaluation** via:
 - Confusion matrix
 - Classification report (precision, recall, F1)
-

Key Findings

- Without class balancing, the model predicted “Pass” almost exclusively.
 - After applying `class_weight='balanced'`, the model identified:
 - ~56% of failing students (recall)
 - ~64% overall accuracy
 - **Demographics alone are weak predictors**, suggesting the need for behavioral/academic data to improve model performance.
 - This helps, though, to support the fairness of academic work across demographics. We can infer that there does not seem to be disparate impact.
-

Next Steps

- Add behavioral predictors like:
 - Assessment scores
 - VLE interaction time
- Compare logistic regression to:
 - Decision trees
 - Random forests
- Tune thresholds and visualize precision-recall tradeoffs

```
import sqlite3
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from ipywidgets import interact, Dropdown

sns.set(style="whitegrid", palette="muted", font_scale=1.1)

conn = sqlite3.connect("open_university.db")

modules = pd.read_sql_query("SELECT DISTINCT code_module FROM studentInfo;", conn)
modules_list = sorted(modules['code_module'].tolist())
modules_list

↪ ['AAA', 'BBB', 'CCC', 'DDD', 'EEE', 'FFF', 'GGG']

def plot_module_scores(selected_module):
    query = f"""
    SELECT
        si.age_band,
        ROUND(AVG(sa.score), 2) AS avg_score
```

```

FROM studentInfo si
JOIN studentAssessment sa ON si.id_student = sa.id_student
WHERE si.code_module = '{selected_module}'
GROUP BY si.age_band
ORDER BY si.age_band;
"""

```

```
df = pd.read_sql_query(query, conn)
```

```

plt.figure(figsize=(8, 5))
sns.barplot(
    data=df,
    x="age_band",
    y="avg_score",
    hue="age_band",          # Assign hue to suppress warning
    palette="crest",
    legend=False             # Hide redundant legend
)
plt.ylim(0, 100)
plt.xticks(rotation=0)
plt.tight_layout()
plt.show()

```

```
interact(plot_module_scores, selected_module=Dropdown(options=modules_list, description="Module:"))
```

```

↪ interactive(children=(Dropdown(description='Module:', options=('AAA', 'BBB', 'CCC',
    'DDD', 'EEE', 'FFF', 'GGG')

```

```
import statsmodels.api as sm
```

```

query = """
SELECT
    si.gender,
    si.age_band,
    AVG(sa.score) AS avg_score
FROM studentInfo si
JOIN studentAssessment sa ON si.id_student = sa.id_student
WHERE sa.score IS NOT NULL
GROUP BY si.id_student;
"""

```

```

df = pd.read_sql_query(query, conn)
conn.close()

```

```

# One-hot encode gender and age_band
df_encoded = pd.get_dummies(df, columns=["gender", "age_band"], drop_first=True)

```

```

# Features and outcome
X = df_encoded.drop("avg_score", axis=1)

```

```

y = df_encoded["avg_score"]
X = sm.add_constant(X)

X = X.astype(float)
y = y.astype(float)

# Fit model
model = sm.OLS(y, X).fit()
print(model.summary())

```



OLS Regression Results

Dep. Variable:	avg_score	R-squared:	0.005			
Model:	OLS	Adj. R-squared:	0.005			
Method:	Least Squares	F-statistic:	40.28			
Date:	Wed, 18 Jun 2025	Prob (F-statistic):	5.91e-26			
Time:	10:26:29	Log-Likelihood:	-97136.			
No. Observations:	23351	AIC:	1.943e+05			
Df Residuals:	23347	BIC:	1.943e+05			
Df Model:	3					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	72.3368	0.164	440.347	0.000	72.015	72.659
gender_M	0.1385	0.204	0.680	0.497	-0.261	0.538
age_band_35-55	2.3370	0.222	10.535	0.000	1.902	2.772
age_band_55<=	4.5241	1.237	3.657	0.000	2.099	6.949
=====						
Omnibus:	5110.201	Durbin-Watson:	1.988			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	11815.804			
Skew:	-1.238	Prob(JB):	0.00			
Kurtosis:	5.453	Cond. No.	14.6			
=====						

Notes:

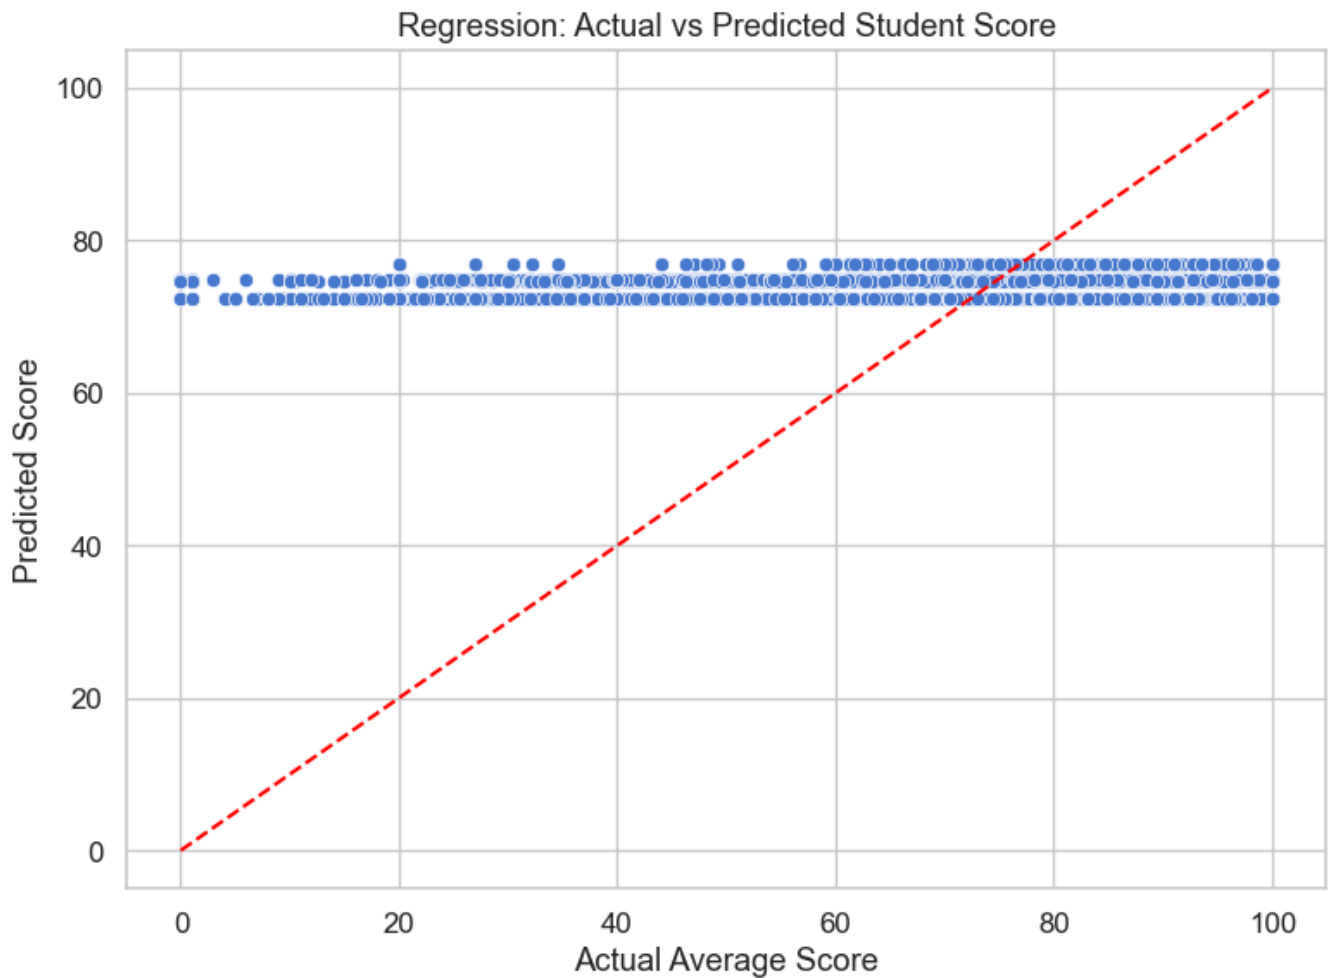
[1] Standard Errors assume that the covariance matrix of the errors is correctly specifi



```
df_encoded["predicted"] = model.predict(X)
```

```
plt.figure(figsize=(8, 6))
sns.scatterplot(x=df_encoded["avg_score"], y=df_encoded["predicted"])
plt.plot([df_encoded["avg_score"].min(), df_encoded["avg_score"].max()],
         [df_encoded["avg_score"].min(), df_encoded["avg_score"].max()],
         color="red", linestyle="--")
```

```
plt.xlabel("Actual Average Score")
plt.ylabel("Predicted Score")
plt.title("Regression: Actual vs Predicted Student Score")
plt.grid(True)
plt.tight_layout()
plt.show()
```



```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, classification_report, ConfusionMatrixDisplay
from sklearn.model_selection import train_test_split
```

```

conn = sqlite3.connect("open_university.db")

query = """
SELECT
    si.gender,
    si.age_band,
    si.region,
    si.disability,
    si.final_result
FROM studentInfo si
WHERE si.final_result IN ('Pass', 'Fail')
"""

df = pd.read_sql_query(query, conn)
conn.close()

df['final_result'] = df['final_result'].map({'Pass': 1, 'Fail': 0})

df_encoded = pd.get_dummies(df, columns=["gender", "age_band", "region", "disability"], drop

X = df_encoded.drop("final_result", axis=1)
y = df_encoded["final_result"]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)

model = LogisticRegression(class_weight='balanced', max_iter=1000)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)

cm = confusion_matrix(y_test, y_pred)
report = classification_report(y_test, y_pred)

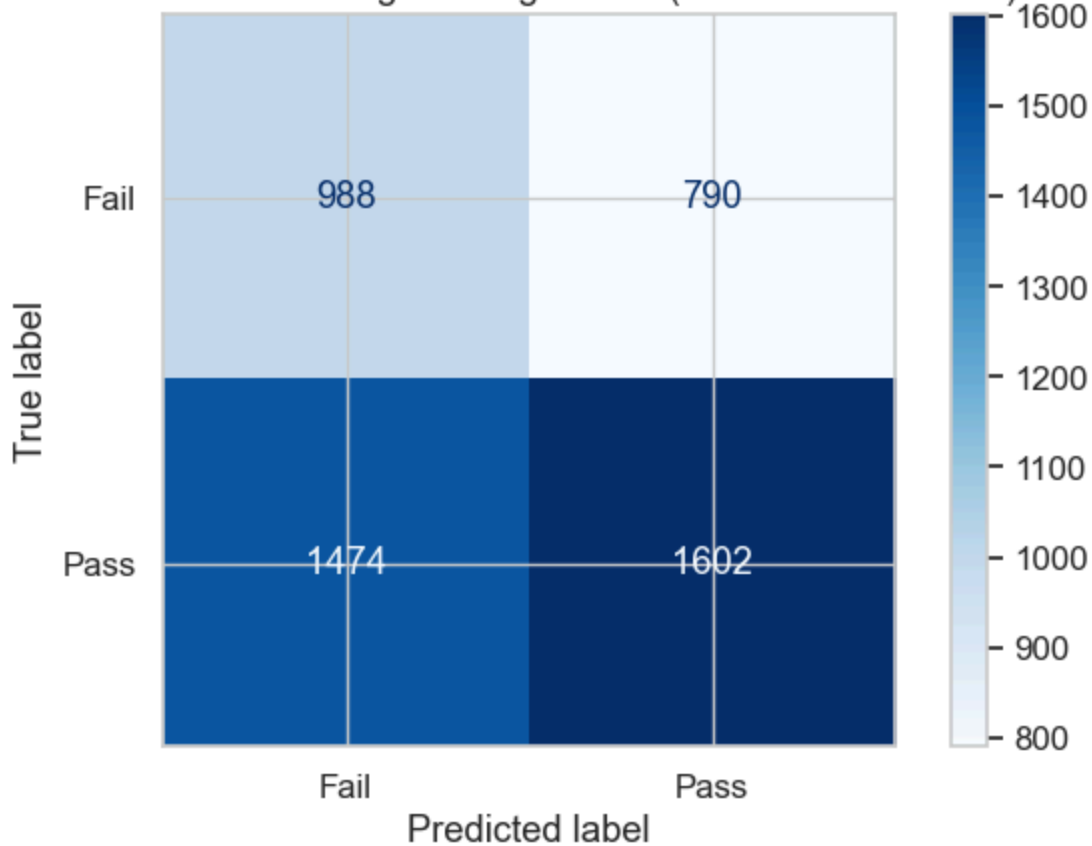
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=["Fail", "Pass"])
disp.plot(cmap="Blues")
plt.title("Confusion Matrix: Logistic Regression (Pass/Fail Prediction)")
plt.tight_layout()
plt.show()

report

```



Confusion Matrix: Logistic Regression (Pass/Fail Prediction)



	precision	recall	f1-score	support
0.56	0.47	1778	1	0.67
		0.52		0.59
				0.40
				3076